

# Study Material - Gen AI Week 2

## Document Information

- **Generated:** 2025-08-02 13:45:00
- **Source:** Handwritten lecture slides
- **Platform:** IITM Course Material
- **Word Count:** ~3,500 words
- **Estimated Reading Time:** ~18 minutes
- **Number of Chapters:** 6
- **Content:** Variational Divergence Minimization (VDM) and Generative Adversarial Networks (GANs)

## Table of Contents

1. Realization of Variational Divergence Minimization (VDM)
  2. Implementing VDM for Generative Modeling
  3. Generative Adversarial Networks (GANs)
  4. GAN Architecture and Implementation
  5. Training GANs: Discriminator and Generator
  6. Practical Implementation of GAN Training
- 

## Lecture Overview

This week's material focuses on the practical realization of Variational Divergence Minimization (VDM) and its application in Generative Adversarial Networks (GANs). The lectures provide a mathematical foundation for understanding how GANs work, starting from the theoretical framework of f-divergences and leading to the practical implementation details of training adversarial networks.

## Learning Objectives

Upon completing this material, students will be able to: - Understand the mathematical foundation of Variational Divergence Minimization (VDM) - Implement VDM using neural networks for generative modeling - Comprehend the architecture and training procedure of Generative Adversarial Networks (GANs) - Analyze the min-max optimization problem in adversarial training - Implement practical GAN training algorithms with alternating optimization - Understand the role of discriminator and generator networks in the adversarial framework

## Prerequisites

To fully grasp the concepts in this material, students should be familiar with: - **Deep Learning Fundamentals:** Neural network architectures, backpropagation, gradient descent - **Probability Theory:** Probability distributions, expectation, divergences (KL divergence, JS divergence) - **Information Theory:** f-divergences, mutual information concepts - **Optimization Theory:** Min-max optimization, saddle point problems - **Mathematical Analysis:** Convex optimization, duality theory

---

## Realization of Variational Divergence Minimization (VDM)

### Mathematical Foundation of VDM

Variational Divergence Minimization provides a framework for training generative models by minimizing the divergence between the data distribution and the model distribution.

## Problem Setup

Given data  $D = \{x_1, x_2, \dots, x_n\}$  drawn i.i.d. from  $p_x$ , we want to learn a generator  $g_\theta(z)$  where  $z \sim \mathcal{N}(0, I)$  such that  $\hat{x} \sim p_\theta(\cdot)$  approximates the true data distribution.

The objective is to find:

$$\theta^* = \arg \min_{\theta} D_f(p_x \| p_\theta)$$

where  $D_f$  represents an f-divergence between the true data distribution  $p_x$  and the model distribution  $p_\theta$ .

## f-Divergence and Variational Representation

For any f-divergence, we can use the variational representation:

$$D_f(p_x \| p_\theta) \geq \max_{T(x) \in \mathcal{T}} [\mathbb{E}_{p_x}[T(x)] - \mathbb{E}_{p_\theta}[f^*(T(x))]]$$

where: -  $f^*$  is the convex conjugate of the f-divergence function  $f$  -  $\mathcal{T}$  is a class of functions (typically neural networks) -  $T(x)$  is the critic function

## Neural Network Parameterization

The optimal solution becomes:

$$\begin{aligned} \theta^* &= \arg \min_{\theta} D_f(p_x \| p_\theta) \\ &\approx \arg \min_{\theta} [\text{lower bound on } D_f] \\ &= \arg \min_{\theta} \max_{T(x)} (\mathbb{E}_{p_x}[T(x)] - \mathbb{E}_{p_\theta}[f^*(T(x))]) \end{aligned}$$

We represent the critic function  $T$  via neural networks  $T_w(x)$  where  $w$  are the parameters of the network.

With this parameterization, the objective becomes:

$$\theta^*, w^* = \arg \min_{\theta} \max_w [\mathbb{E}_{p_x}[T_w(x)] - \mathbb{E}_{p_\theta}[f^*(T_w(x))]]$$

---

# Implementing VDM for Generative Modeling

## Architecture Overview

The VDM framework consists of two main components:

### 1. Generator Network

- **Input:** Random noise  $z \sim \mathcal{N}(0, I)$
- **Function:**  $g_\theta(z) \rightarrow \hat{x} \sim p_\theta(\cdot)$
- **Purpose:** Generate synthetic data samples

### 2. Critic Network (Discriminator)

- **Input:** Data samples  $x$  (real or generated)
- **Function:**  $T_w(x) \rightarrow T_w(x)$
- **Purpose:** Evaluate the quality of samples according to the f-divergence

## Mathematical Formulation

The complete objective function becomes:

$$J(\theta, w) = \mathbb{E}_{p_x}[T_w(x)] - \mathbb{E}_{p_\theta}[f^*(T_w(x))]$$

The optimization problem is a saddle point:

$$\theta^*, w^* = \arg \min_{\theta} \max_w J(\theta, w)$$

This is an **adversarial problem** where: - The generator (parameterized by  $\theta$ ) tries to minimize the objective  
- The critic (parameterized by  $w$ ) tries to maximize the objective

## f-Divergence Specific Activations

For different f-divergences, we need specific activation functions:

### General Form

$$T_w(x) = \sigma_f(V_w(x))$$

where: -  $V_w(x) : \mathcal{X} \rightarrow \mathbb{R}$  is a neural network -  $\sigma_f(v) : \mathbb{R} \rightarrow \text{dom } f^*$  is the f-divergence specific activation -  $\text{dom } f^*$  is the domain of the conjugate function  $f^*$

The final critic output:  $T_w(x) : \mathcal{X} \rightarrow \text{dom } f^*$

---

## Generative Adversarial Networks (GANs)

### GAN as a Specific Case of VDM

GANs represent a specific instantiation of the VDM framework using the Jensen-Shannon divergence.

### f-Divergence for GANs

For GANs, the f-divergence is defined as:

$$f(u) = u \log u - (u + 1) \log(u + 1)$$

This is similar to the Jensen-Shannon Divergence (JSD).

### Conjugate Function and Activation

The conjugate function is:

$$f^*(t) = -\log(1 - \exp(t)), \quad \text{dom } f^* = \mathbb{R}$$

The activation function becomes:

$$\sigma_f(v) = -\log(1 + e^{-v})$$

## GAN Objective Function

### General VDM Form

$$J(\theta, w) = \mathbb{E}_{p_x}[\sigma_f(V_w(x))] - \mathbb{E}_{p_\theta}[f^*(\sigma_f(V_w(x)))]$$

## GAN-Specific Form


$$J_{GAN}(\theta, w) = \mathbb{E}_{p_x}[\log D_w(x)] + \mathbb{E}_{p_\theta}[\log(1 - D_w(x))]$$

where the discriminator function is defined as:

$$D_w(x) = \frac{1}{1 + e^{-V_w(x)}}$$

This is the **sigmoid function**, which outputs values in  $[0, 1]$ .

## Architecture Diagram



The discriminator acts as a **binary classifier** distinguishing between real and generated samples.

---

# GAN Architecture and Implementation

## Network Architecture

### Generator Network

- **Input:**  $z \sim \mathcal{N}(0, I)$  (random noise)
- **Output:**  $\hat{x} \sim p_{\theta}(x)$  (synthetic data)
- **Architecture:** Multi-layer neural networks (MLP, FNN, CNN)

### Discriminator Network

- **Input:**  $x$  (real or generated samples)
- **Output:**  $D_w(x) \in [0, 1]$  (probability that input is real)
- **Architecture:** Neural networks (CNN for images, MLP for other data)
- **Final layer:** Sigmoid activation for binary classification

## GAN Loss Function

The complete GAN objective is:

$$J_{GAN}(\theta, w) = \mathbb{E}_{x \sim p_x}[\log D_w(x)] + \mathbb{E}_{\hat{x} \sim p_{\theta}}[\log(1 - D_w(\hat{x}))]$$

## Implementation in Practice

### Input Data

Given dataset:  $D = \{x_1, x_2, x_3, \dots, x_n\}$  drawn i.i.d. from  $p_x$

### Discriminator Optimization

$$w^* = \arg \max_w (\mathbb{E}_{p_x}[\log D_w(x)] + \mathbb{E}_{p_{\theta}}[\log(1 - D_w(x))])$$

Using mini-batches  $B_1$  and  $B_2$ :

$$w^* \approx \arg \max_w \left[ \frac{1}{B_1} \sum_{i=1}^{B_1} \log D_w(x_i) + \frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(\hat{x}_j)) \right]$$

where: -  $x_1, \dots, x_{B_1} \sim p_x$  (real samples) -  $\hat{x}_1, \dots, \hat{x}_{B_2} \sim p_\theta$  (generated samples) -  $\hat{x}_j = g_\theta(z_j)$  with  $z_j$  sampled from noise distribution

---

## Training GANs: Discriminator and Generator

### Alternating Optimization Procedure

GAN training involves alternating between optimizing the discriminator and generator:

#### Discriminator Update Step

**Objective:** Maximize discriminator's ability to distinguish real from fake

$$w^{t+1} \leftarrow w^t + \alpha_1 \nabla_w J_{GAN}(\theta, w)$$

**Keep  $\theta$  constant and optimize:**

$$\theta^* = \arg \min_{\theta} J_{GAN}(\theta, w)$$

This simplifies to:

$$\theta^* \approx \arg \min_{\theta} \left[ \frac{1}{B_1} \sum_{i=1}^{B_1} \log D_w(x_i) + \frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j))) \right]$$

Since the first term is independent of  $\theta$ :

$$\theta^* \approx \arg \min_{\theta} \left[ \frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j))) \right]$$

#### Generator Update Step

**Objective:** Minimize generator's loss (fool the discriminator)

$$\theta^{t+1} \leftarrow \theta^t - \alpha_2 \nabla_{\theta} J_{GAN}(\theta, w)$$

**Keep  $w$  constant and optimize:**

$$J_{GAN} = -\frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j)))$$

**Update:**  $\theta^{t+1} \leftarrow \theta^t - \alpha_2 \nabla_{\theta} J_{GAN}(\theta, w)$

---

## Practical Implementation of GAN Training

### Training the Discriminator

#### Step-by-Step Process

1. **Keep  $\theta$  constant** - Fix generator parameters
2. **Sample data:**  $D = \{x_1, x_2, \dots, x_n\}$

### 3. Forward Propagation:

- Real samples:  $x_1, x_2, \dots, x_{B_1} \rightarrow D_w(\cdot) \rightarrow$  classification scores
- Generated samples:  $z_1 \dots z_{B_2} \rightarrow g_\theta(z) \rightarrow g_\theta(z_1) \dots g_\theta(z_{B_2}) \rightarrow D_w(\cdot) \rightarrow$  classification scores

### 4. Compute Loss:

$$J_{GAN}(\theta, w) = \left[ \frac{1}{B_1} \sum_{i=1}^{B_1} \log D_w(x_i) + \frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j))) \right]$$

### 5. Backward Propagation: Compute $\nabla_w J_{GAN}(\theta, w)$

### 6. Update: $w^{t+1} \leftarrow w^t + \alpha_1 \nabla_w J_{GAN}$

## Training the Generator

### Step-by-Step Process

#### 1. Sample noise: $z_1 \dots z_{B_2} \sim \mathcal{N}(0, I)$

#### 2. Forward Propagation:

- $z \rightarrow g_\theta(z) \rightarrow g_\theta(z_1) \dots g_\theta(z_{B_2}) \sim p_\theta$
- Generated samples through discriminator:  $g_\theta(z_j) \rightarrow D_w(\cdot)$

#### 3. Compute Generator Loss:

$$J_{GAN} = -\frac{1}{B_2} \sum_{j=1}^{B_2} \log(1 - D_w(g_\theta(z_j)))$$

#### 4. Update $\theta$ only with $w$ constant:

$$\theta^{t+1} \leftarrow \theta^t - \alpha_2 \nabla_\theta J_{GAN}(\theta, w)$$

## Training Algorithm Summary

1. Initialize  $\theta, w$
2. For epoch = 1 to max\_epochs:
  3. // Train Discriminator
  4. Sample  $\{x, \dots, x_B\}$  from real data
  5. Sample  $\{z, \dots, z_B\} \sim \mathcal{N}(0, I)$
  6. Generate  $\{\hat{x}, \dots, \hat{x}_B\} = \{g(z), \dots, g(z_B)\}$
  7. Compute discriminator loss  $J_D$
  8.  $w^{t+1} \leftarrow w^t + \alpha_1 J_D$
  9. // Train Generator
  10. Sample  $\{z, \dots, z_B\} \sim \mathcal{N}(0, I)$
  11. Compute generator loss  $J_G$
  12.  $\theta^{t+1} \leftarrow \theta^t - \alpha_2 J_G$

## Key Training Insights

### Saddle Point Optimization

- GANs solve a **min-max** problem:  $\min_\theta \max_w J(\theta, w)$
- This is fundamentally different from standard neural network training
- Requires careful balance between discriminator and generator updates

## Training Stability

- **Discriminator too strong:** Generator cannot learn (gradients vanish)
- **Generator too strong:** Discriminator cannot provide useful gradients
- **Balanced training:** Both networks improve together

## Implementation Considerations

- **Alternating updates:** Update one network while keeping the other fixed
  - **Learning rates:** Often different learning rates for generator and discriminator
  - **Update frequency:** Sometimes update discriminator multiple times per generator update
  - **Batch sizes:** Can use different batch sizes for real and generated samples
- 

## Key Takeaways from Week 2

### Theoretical Understanding

1. **VDM Framework:** Provides mathematical foundation for adversarial training
2. **f-Divergences:** Allow different types of distance measures between distributions
3. **Variational Representation:** Enables practical optimization using neural networks

### Practical Implementation

1. **GAN Architecture:** Two-network adversarial system (generator + discriminator)
2. **Training Procedure:** Alternating optimization between networks
3. **Loss Functions:** Binary cross-entropy for discrimination task

### Mathematical Insights

1. **Saddle Point Problem:** Min-max optimization requires special consideration
2. **Neural Network Parameterization:** Both generator and critic are neural networks
3. **Gradient-Based Training:** Standard backpropagation with careful update procedures

### Applications and Extensions

1. **Image Generation:** GANs excel at generating realistic images
2. **Data Augmentation:** Generate synthetic training data
3. **Domain Transfer:** Learn mappings between different data domains

This material provides the foundation for understanding more advanced GAN variants and training techniques that will be covered in subsequent weeks.