

# Study Material - Youtube

## Document Information

- **Generated:** 2025-08-26 06:01:32
- **Source:** <https://www.youtube.com/watch?v=eH9skKyqjJM>
- **Platform:** Youtube
- **Word Count:** 2,320 words
- **Estimated Reading Time:** ~11 minutes
- **Number of Chapters:** 25
- **Transcript Available:** Yes (analyzed from video content)

## Table of Contents

1. Bidirectional GANs (Bi-GANs): A Deep Dive
2. Imports
3. Configuration
4. At 9:05, the instructor defines the Generator
5. At 9:37, the instructor defines the Encoder
6. At 9:57, the instructor defines the Discriminator
7. Concatenate image and latent vector
8. Optimizer for the Discriminator
9. Joint optimizer for the Generator and Encoder
10. — Train Discriminator —
11. Set models to correct modes
12. 1. Real pairs: (real\_image, encoded\_latent\_vector)
13. 2. Fake pairs: (generated\_image, original\_latent\_vector)
14. Calculate loss and update weights
15. — Train Generator + Encoder —
16. Set models to correct modes
17. Generate fake image and fake latent vector
18. Get discriminator scores
19. Calculate loss to fool the discriminator
20. Note: We want the discriminator to output 1 for both pairs
21. A more intuitive way is:
22.  $\text{loss\_GE} = \text{bce\_loss}(D\_real, \text{torch.ones\_like}(D\_real)) + \text{bce\_loss}(D\_fake, \text{torch.ones\_like}(D\_fake))$
23. Update weights
24. Self-Assessment for This Video
25. Key Takeaways from This Video

---

## Video Overview

This video provides a comprehensive tutorial on **Bidirectional Generative Adversarial Networks (Bi-GANs)**. It begins by establishing the foundational concepts of a standard “vanilla” GAN, highlighting its limitation in mapping from the data space back to the latent space. The lecture then introduces the Bi-GAN architecture as a solution, detailing its three core components: a Generator, an Encoder, and a joint Discriminator. The instructor meticulously explains the mathematical mappings and the information flow for both real and generated data pairs. The second half of the tutorial is a practical, hands-on coding session in a Google Colab environment, demonstrating how to implement a Bi-GAN from scratch using PyTorch to train on the MNIST dataset.

## Learning Objectives

Upon completing this study material, you will be able to: - **Understand the architecture of a standard GAN** and its primary function. - **Identify the key limitation of a standard GAN**, namely the absence of an inverse mapping from data to latent space. - **Describe the architecture of a Bidirectional GAN (Bi-GAN)** and its three components: Generator, Encoder, and Discriminator. - **Explain the mathematical mappings** for each component in the Bi-GAN framework. - **Understand the role of the joint discriminator** in evaluating pairs of images and latent vectors. - **Formulate the Bi-GAN loss function** and understand the adversarial training dynamics between the components. - **Implement a Bi-GAN in PyTorch**, including defining the network architectures, setting up optimizers, and structuring the training loop.

## Prerequisites

To fully grasp the concepts in this video, you should have a solid understanding of: - **Fundamentals of Generative Adversarial Networks (GANs)**. - **Neural Network principles**, including linear layers, activation functions (ReLU, Sigmoid, Tanh), and batch normalization. - **Basic probability and calculus**, particularly the concepts of probability distributions and expectation. - **Python programming** and familiarity with the **PyTorch** deep learning library.

## Key Concepts Covered

- Standard (Vanilla) GAN
  - Bidirectional GAN (Bi-GAN)
  - Generator ( $G_\theta$ )
  - Encoder ( $E_\phi$ )
  - Joint Discriminator ( $D_\omega$ )
  - Latent Space ( $Z$ ) and Data Space ( $X$ )
  - Latent Inversion / Embedding
  - Bi-GAN Loss Function
  - PyTorch Implementation (`nn.Module`, `nn.Sequential`, `optim.Adam`)
  - Adversarial Training Loop
- 

# Bidirectional GANs (Bi-GANs): A Deep Dive

## From Unidirectional to Bidirectional: The Motivation

A standard Generative Adversarial Network (GAN) is a powerful tool for learning a data distribution and generating new samples from it. However, its information flow is primarily **unidirectional**. It provides a mapping from a simple latent space (e.g., a Gaussian distribution) to the complex data space (e.g., images), but it does not inherently learn the reverse mapping.

As the instructor points out at (1:39), a standard GAN lacks a mechanism for **inversion**—that is, taking a real image from the data space and finding its corresponding representation or “embedding” in the latent space. This capability is useful for a variety of tasks, such as feature extraction and image manipulation.

**Bi-GANs** were introduced to solve this problem by simultaneously learning both the forward (generation) and reverse (encoding) mappings, creating a “bidirectional” flow of information.

## Standard GAN Architecture: A Recap

Before diving into Bi-GANs, the instructor provides a quick review of the standard GAN formulation (00:47). It consists of two networks:

1. **The Generator ( $G_\theta$ ):** This network learns to map a point from a simple prior distribution (latent space  $Z$ ) to the data distribution (image space  $X$ ).

- **Intuition:** The generator is like an artist who learns to create realistic paintings (images) from a simple idea or inspiration (a latent vector).
- **Mathematical Mapping:**

$$G_\theta : Z \rightarrow X$$

Where  $z \in Z$  is a latent vector sampled from a prior distribution  $p_z$  (e.g.,  $\mathcal{N}(0, I)$ ), and  $G_\theta(z)$  is the generated image.

2. **The Discriminator ( $D_\omega$ ):** This network is a binary classifier that learns to distinguish between real data samples from  $X$  and fake samples generated by  $G_\theta$ .

- **Intuition:** The discriminator is like an art critic who learns to tell the difference between a real masterpiece and a forgery.
- **Mathematical Mapping:**

$$D_\omega : X \rightarrow [0, 1]$$

Where  $D_\omega(x)$  represents the probability that the input image  $x$  is real.

The information flow in a standard GAN can be visualized as follows:

graph TD

```

subgraph Generator Path
    Z["Latent Vector (z)"] --> G["Generator (G)"];
    G --> X_fake["Fake Image (x̂)"];
end
subgraph Real Data Path
    X_real["Real Image (x)"] --> D["Discriminator (D)"];
end
X_fake --> D;
D --> Score["Real/Fake Score"];

```

Figure 1: Information flow in a standard GAN. The generator creates fake images, and the discriminator evaluates both real and fake images.

## The Bi-GAN Architecture

The Bi-GAN architecture (1:57) extends the standard GAN by adding a third component, the **Encoder**, and modifying the role of the discriminator.

### The Three Key Players: Generator, Encoder, and Discriminator

1. **The Generator ( $G_\theta$ ):** Its role remains the same as in a standard GAN. It maps latent vectors to images.

$$G_\theta : Z \rightarrow X$$

2. **The Encoder ( $E_\phi$ ):** This is the new component that provides the reverse mapping. It takes a real image from the data space  $X$  and maps it to a representation in the latent space  $Z$ .

- **Intuition:** The encoder acts as an “inverse generator.” It looks at a real image and tries to figure out the latent code that could have produced it.
- **Mathematical Mapping (2:09):**

$$E_\phi : X \rightarrow Z$$

3. **The Joint Discriminator ( $D_\omega$ ):** The discriminator in a Bi-GAN is fundamentally different. Instead of evaluating just an image, it evaluates a **joint pair** of an image and a latent vector  $(x, z)$ .

- **Intuition:** The discriminator now acts as a “consistency judge.” It checks if an image and its corresponding latent code are a plausible pair. It learns to distinguish between pairs that come

from the “real” data flow (a real image and its encoded representation) and pairs that come from the “fake” data flow (a generated image and its original latent code).

- **Mathematical Mapping (2:18):**

$$D_\omega : X \times Z \rightarrow [0, 1]$$

## Visualizing the Information Flow in Bi-GAN

The Bi-GAN framework involves two parallel information streams that feed into the discriminator, as illustrated by the instructor’s drawing (2:56) and visualized below.

graph TD

```

subgraph Generator Path (Fake Pair)
    Z_prior["Latent Vector z ~ p(z)"] --> G["Generator G_"];
    G --> X_fake["Fake Image x̂ = G(z)"];
    subgraph Fake Pair
        X_fake --> D["Discriminator D_"];
        Z_prior --> D;
    end
end

subgraph Encoder Path (Real Pair)
    X_real["Real Image x ~ p_data(x)"] --> E["Encoder E_"];
    E --> Z_fake["Encoded Vector ẑ = E(x)"];
    subgraph Real Pair
        X_real --> D;
        Z_fake --> D;
    end
end

D --> Score["Real/Fake Pair Score"];

```

Figure 2: The bidirectional information flow in a Bi-GAN. The discriminator learns to distinguish between real pairs (real image, its encoding) and fake pairs (generated image, its original latent vector).

- **Real Pair:** A real image  $x$  is passed through the encoder  $E_\phi$  to get its latent representation  $\hat{z} = E_\phi(x)$ . The pair  $(x, \hat{z})$  is considered “real.”
- **Fake Pair:** A latent vector  $z$  is sampled from the prior and passed through the generator  $G_\theta$  to create a fake image  $\hat{x} = G_\theta(z)$ . The pair  $(\hat{x}, z)$  is considered “fake.”

The discriminator  $D_\omega$  is trained to output a high probability (close to 1) for real pairs and a low probability (close to 0) for fake pairs. The generator and encoder are jointly trained to produce pairs that fool the discriminator into thinking they are real.

## Key Mathematical Concepts

### The Bi-GAN Objective Function

The training process is a minimax game involving the three networks. The objective function, which the discriminator tries to maximize and the generator/encoder try to minimize, is formulated as follows (5:48):

$$\max_D \min_{G, E} V(D, G, E) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x, E(x))] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z), z))]$$

Let’s break down this equation:

- $\mathbb{E}_{x \sim p_{data}(x)} [\log D(x, E(x))]$ : This is the term for **real pairs**.  
 –  $x \sim p_{data}(x)$ : We sample a real image  $x$  from the true data distribution.

- $E(x)$ : We encode the real image to get its latent representation.
- $D(x, E(x))$ : The discriminator evaluates this “real” pair.
- The discriminator  $D$  is trained to maximize this term (i.e., make  $D(x, E(x))$  close to 1), which means correctly identifying the pair as real.
- $\mathbb{E}_{z \sim p_z(z)}[\log(1 - D(G(z), z))]$ : This is the term for **fake pairs**.
  - $z \sim p_z(z)$ : We sample a random latent vector  $z$  from the prior distribution.
  - $G(z)$ : We generate a fake image from this latent vector.
  - $D(G(z), z)$ : The discriminator evaluates this “fake” pair.
  - The discriminator  $D$  is trained to make  $D(G(z), z)$  close to 0, which maximizes  $\log(1 - D(\cdot))$ . This means correctly identifying the pair as fake.

The generator  $G$  and encoder  $E$  are trained to do the opposite: they want to **minimize** this objective. This forces them to create pairs  $(G(z), z)$  and  $(x, E(x))$  that the discriminator misclassifies as real.

## Practical Implementation in PyTorch

The instructor demonstrates a full implementation of Bi-GAN in a Google Colab notebook (8:37).

### Setup and Configuration

Standard libraries are imported, and hyperparameters are defined.

```
# Imports
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
import numpy as np
import matplotlib.pyplot as plt

# Configuration
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
batch_size = 128
latent_dim = 50
epochs = 30
lr = 2e-4
```

### Model Definitions

Three `nn.Module` classes are defined for the Generator, Encoder, and Discriminator.

**Generator Network** The generator maps the `latent_dim` to the flattened image size (784 for MNIST). The final activation is `nn.Tanh()` (14:19), which scales the output to the range  $[-1, 1]$ , suitable for normalized image data.

```
# At 9:05, the instructor defines the Generator
class Generator(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(latent_dim, 1024),
            nn.ReLU(),
            nn.Linear(1024, 1024),
            nn.BatchNorm1d(1024),
            nn.ReLU(),
```

```

        nn.Linear(1024, 784),
        nn.Tanh() # Corrected from Sigmoid
    )
    def forward(self, z):
        return self.net(z)

```

**Encoder Network** The encoder performs the reverse mapping, from the image space (784) to the latent space (latent\_dim).

*# At 9:37, the instructor defines the Encoder*

```

class Encoder(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(784, 1024),
            nn.ReLU(),
            nn.Linear(1024, 1024),
            nn.BatchNorm1d(1024),
            nn.ReLU(),
            nn.Linear(1024, latent_dim)
        )
    def forward(self, x):
        return self.net(x)

```

**Discriminator Network** The discriminator takes a concatenated input of the image and the latent vector. Its input dimension is 784 + latent\_dim. It outputs a single logit.

*# At 9:57, the instructor defines the Discriminator*

```

class Discriminator(nn.Module):
    def __init__(self):
        super().__init__()
        self.net = nn.Sequential(
            nn.Linear(784 + latent_dim, 1024),
            nn.ReLU(),
            nn.Linear(1024, 1024),
            nn.BatchNorm1d(1024),
            nn.ReLU(),
            nn.Linear(1024, 1) # No sigmoid, as BCEWithLogitsLoss is used
        )
    def forward(self, x, z):
        # Concatenate image and latent vector
        xz = torch.cat((x, z), dim=1)
        return self.net(xz)

```

## Optimizers and Loss Function

Two separate optimizers are created (10:36): one for the discriminator and a joint one for the generator and encoder.

```
bce_loss = nn.BCEWithLogitsLoss()
```

*# Optimizer for the Discriminator*

```
optimizer_D = optim.Adam(D.parameters(), lr=lr, betas=(0.5, 0.999))
```

*# Joint optimizer for the Generator and Encoder*

```
optimizer_GE = optim.Adam(list(G.parameters()) + list(E.parameters()), lr=lr, betas=(0.5, 0.999))
```

## The Training Loop: A Step-by-Step Guide

The training loop alternates between training the discriminator and training the generator/encoder.

**Phase 1: Training the Discriminator (11:13)** The goal is to teach the discriminator to distinguish real pairs from fake pairs.

```
# --- Train Discriminator ---
# Set models to correct modes
G.eval()
E.eval()
D.train()

# 1. Real pairs: (real_image, encoded_latent_vector)
z_fake = E(x_real).detach()
D_real = D(x_real, z_fake)
label_real = torch.ones_like(D_real)

# 2. Fake pairs: (generated_image, original_latent_vector)
x_fake = G(z_real).detach()
D_fake = D(x_fake, z_real)
label_fake = torch.zeros_like(D_fake)

# Calculate loss and update weights
loss_D = bce_loss(D_real, label_real) + bce_loss(D_fake, label_fake)
optimizer_D.zero_grad()
loss_D.backward()
optimizer_D.step()
```

**Phase 2: Training the Generator and Encoder (13:42)** The goal is to train the generator and encoder to fool the discriminator.

```
# --- Train Generator + Encoder ---
# Set models to correct modes
G.train()
E.train()
D.eval()

# Generate fake image and fake latent vector
x_fake = G(z_real)
z_fake = E(x_real)

# Get discriminator scores
D_real = D(x_real, z_fake)
D_fake = D(x_fake, z_real)

# Calculate loss to fool the discriminator
# Note: We want the discriminator to output 1 for both pairs
loss_GE = bce_loss(D_real, label_fake) + bce_loss(D_fake, label_real) # This is a common implementation
# A more intuitive way is:
# loss_GE = bce_loss(D_real, torch.ones_like(D_real)) + bce_loss(D_fake, torch.ones_like(D_fake))

# Update weights
optimizer_GE.zero_grad()
loss_GE.backward()
```

`optimizer_GE.step()`

---

## Self-Assessment for This Video

1. **Conceptual Question:** What is the primary architectural difference between a standard GAN and a Bi-GAN? Why is this difference significant?
  2. **Mathematical Question:** Write down the mathematical mapping for the Bi-GAN discriminator. What are its inputs and what does its output represent?
  3. **Intuition Question:** Explain in your own words the two types of pairs (real and fake) that the Bi-GAN discriminator is trained on.
  4. **Coding Question:** In the PyTorch implementation, why are the generator and encoder parameters updated using a single, combined optimizer (`optimizer_GE`)?
  5. **Application Question:** Why is learning an encoder (the  $X \rightarrow Z$  mapping) a useful capability for a generative model?
- 

## Key Takeaways from This Video

- **Bi-GANs create a bidirectional mapping:** They learn to map from latent space to data space ( $G : Z \rightarrow X$ ) and from data space back to latent space ( $E : X \rightarrow Z$ ).
- **The architecture has three components:** A Generator ( $G$ ), an Encoder ( $E$ ), and a joint Discriminator ( $D$ ).
- **The Discriminator judges pairs:** It learns to distinguish between “real” pairs of (real image, its encoding) and “fake” pairs of (generated image, its original latent code).
- **The loss function reflects the three-player game:** The discriminator is trained to maximize classification accuracy, while the generator and encoder are jointly trained to minimize it, thereby fooling the discriminator.
- **Implementation requires careful management:** The training loop must alternate between updating the discriminator and updating the generator/encoder, with models set to the correct `train()` or `eval()` modes in each phase.