

# Study Material - Youtube

## Document Information

- **Generated:** 2025-08-26 06:13:02
- **Source:** <https://www.youtube.com/watch?v=zUJNypPc-Vo>
- **Platform:** Youtube
- **Word Count:** 2,037 words
- **Estimated Reading Time:** ~10 minutes
- **Number of Chapters:** 6
- **Transcript Available:** Yes (analyzed from video content)

## Table of Contents

1. From Maximum Likelihood to ELBO Maximization
  2. The Expectation-Maximization (EM) Algorithm
  3. Application: Gaussian Mixture Models (GMM)
  4. Limitations and The Path to Deep Generative Models
  5. Self-Assessment for This Video
  6. Key Takeaways from This Video
- 

## Video Overview

This lecture provides a concise yet comprehensive review of the mathematical framework for training latent variable models, culminating in the formulation of the Expectation-Maximization (EM) algorithm. The instructor begins by recapping the derivation of the **Evidence Lower Bound (ELBO)**, which serves as a tractable surrogate for the intractable log-likelihood of the data. The core of the lecture is framing the learning problem as a **joint optimization** of the ELBO with respect to both the model parameters ( $\theta$ ) and a variational distribution ( $q$ ). This general framework is then illustrated with a classic machine learning example, the **Gaussian Mixture Model (GMM)**, setting the stage for understanding how these foundational concepts are extended to more complex deep generative models.

## Learning Objectives

Upon completing this lecture, students will be able to: - **Recall** the derivation of the Evidence Lower Bound (ELBO) from the log-likelihood using Jensen's Inequality. - **Formulate** the joint optimization problem for latent variable models, which involves maximizing the ELBO with respect to both model parameters ( $\theta$ ) and the variational posterior ( $q$ ). - **Understand** the iterative nature of the Expectation-Maximization (EM) algorithm as a method to solve this joint optimization. - **Define** a Gaussian Mixture Model (GMM) as a specific instance of a latent variable model with discrete latent variables. - **Identify** the key parameters of a GMM that need to be learned from data. - **Recognize** the limitations of the standard EM algorithm when the true posterior is intractable, motivating the need for more advanced techniques like those used in VAEs.

## Prerequisites

To fully grasp the concepts in this lecture, students should have a solid understanding of: - **Probability Theory:** Concepts of joint, conditional, and marginal probability distributions, Bayes' theorem, and the definition of expectation. - **Calculus:** Multivariate differentiation and optimization (finding maxima by setting derivatives to zero). - **Machine Learning Fundamentals:** Basic principles of Maximum Likelihood Estimation (MLE) and the concept of latent variables. - **Information Theory:** A basic understanding of Jensen's Inequality is crucial for the ELBO derivation.

## Key Concepts Covered in This Video

- Latent Variable Models
  - Log-Likelihood Maximization
  - Evidence Lower Bound (ELBO)
  - Jensen's Inequality
  - Variational Inference
  - Joint Optimization Problem
  - Expectation-Maximization (EM) Algorithm
  - Gaussian Mixture Models (GMM)
- 

## From Maximum Likelihood to ELBO Maximization

This section recaps the foundational challenge in training latent variable models and how the Evidence Lower Bound (ELBO) provides a tractable solution.

### The Challenge of Latent Variable Models

(00:15) The primary goal in many generative models is to perform **Maximum Likelihood Estimation (MLE)**. We want to find the model parameters  $\theta$  that maximize the likelihood of our observed data  $X$ . This is typically done by maximizing the log-likelihood.

For a latent variable model, the probability of an observed data point  $x$ ,  $p_\theta(x)$ , is obtained by marginalizing over the unobserved (latent) variable  $z$ :

$$p_\theta(x) = \int_z p_\theta(x, z) dz$$

The optimization problem is to find the optimal parameters  $\theta^*$ :

$$\theta^* = \arg \max_{\theta} \mathbb{E}_{p_x} [\log p_\theta(x)]$$

The core difficulty arises because the log-likelihood function,  $L(\theta) = \log p_\theta(x)$ , involves a logarithm of an integral (or a sum for discrete latent variables):

$$L(\theta) = \log \int_z p_\theta(x, z) dz$$

This expression is often **intractable**, meaning it cannot be computed or optimized analytically. The logarithm prevents us from pushing the integral outside, making direct optimization impossible for most non-trivial models.

### Deriving the Evidence Lower Bound (ELBO)

To overcome the intractability of the log-likelihood, we introduce a tractable lower bound. This is achieved through **variational inference**, where we introduce an auxiliary distribution  $q(z|x)$  to approximate the true, but unknown, posterior distribution of the latent variables,  $p_\theta(z|x)$ .

The derivation of the ELBO, as recapped by the instructor (00:38), proceeds as follows:

1. **Start with the log-likelihood** and introduce  $q(z|x)$  by multiplying and dividing inside the integral:

$$L(\theta) = \log \int_z p_\theta(x, z) dz = \log \int_z q(z|x) \frac{p_\theta(x, z)}{q(z|x)} dz$$

2. **Recognize the integral as an expectation.** The expression is now the logarithm of an expectation with respect to the distribution  $q(z|x)$ :

$$L(\theta) = \log \mathbb{E}_{q(z|x)} \left[ \frac{p_\theta(x, z)}{q(z|x)} \right]$$

3. **Apply Jensen's Inequality.** (01:20) For a concave function like the logarithm, the log of an expectation is greater than or equal to the expectation of the log:  $\log \mathbb{E}[Y] \geq \mathbb{E}[\log Y]$ . Applying this gives us:

$$L(\theta) \geq \mathbb{E}_{q(z|x)} \left[ \log \left( \frac{p_\theta(x, z)}{q(z|x)} \right) \right]$$

4. **Define the ELBO.** This lower bound is known as the **Evidence Lower Bound (ELBO)**, which we denote as  $J_\theta(q)$ :

$$J_\theta(q) = \mathbb{E}_{q(z|x)} [\log p_\theta(x, z) - \log q(z|x)]$$

Thus, we have the fundamental relationship:

$$L(\theta) \geq J_\theta(q)$$

The following flowchart, inspired by the instructor's recap, illustrates this derivation process.

flowchart TD

```

A["Start: Log-Likelihood<br/>L() = log p(x)"] --> B["Introduce Latent Variable z<br/>L() = log p(x, z) - log q(z|x)"]
B --> C["Introduce Variational Posterior q(z|x)<br/>L() = log \int q(z|x) [p(x, z)/q(z|x)] dz"]
C --> D["Recognize as Log of Expectation<br/>L() = log E_q[p(x, z)/q(z|x)]"]
D --> E["Apply Jensen's Inequality<br/>log E[Y] \ge E[log Y]"]
E --> F["Obtain Evidence Lower Bound (ELBO)<br/>L() \ge E_q[log p(x, z) - log q(z|x)]"]
F --> G["New Goal: Maximize ELBO<br/>max_{\theta, q} J(\theta, q)"]

```

*This flowchart shows the key steps in deriving the Evidence Lower Bound (ELBO) as a tractable objective function for latent variable models.*

## The Joint Optimization Problem

(01:51) The ELBO is a function of both the model parameters  $\theta$  and the variational distribution  $q$ . Therefore, our new goal is to **jointly optimize** the ELBO with respect to both:

$$\theta^*, q^* = \arg \max_{\theta, q} J_\theta(q) = \arg \max_{\theta, q} \mathbb{E}_{q(z|x)} \left[ \log \frac{p_\theta(x, z)}{q(z|x)} \right]$$

This is the fundamental optimization problem that is solved in any latent variable generative model, including VAEs and Diffusion Models.

## The Expectation-Maximization (EM) Algorithm

The joint optimization problem is typically solved with an iterative procedure known as the **Expectation-Maximization (EM) algorithm**. This algorithm breaks the difficult joint optimization into two simpler, alternating steps.

## The Iterative EM Procedure

(13:07) The EM algorithm iteratively refines the estimates for  $\theta$  and  $q$ . Let  $\theta_t$  and  $q_t$  be the estimates at iteration  $t$ . The algorithm proceeds as follows:

1. **E-Step (Expectation Step):** (14:22) Keep the model parameters  $\theta_t$  fixed and find the optimal distribution  $q_{t+1}^*$  that maximizes the ELBO.

$$q_{t+1}^* = \arg \max_q J_{\theta_t}(q)$$

It can be analytically shown that the optimal  $q$  that maximizes this bound is the true posterior distribution of the latent variables, conditioned on the data and the current parameters:


$$q_{t+1}^*(z|x) = p_{\theta_t}(z|x)$$

2. **M-Step (Maximization Step):** (16:24) Keep the distribution  $q_{t+1}^*$  fixed and find the new model parameters  $\theta_{t+1}^*$  that maximize the ELBO.

$$\theta_{t+1}^* = \arg \max_{\theta} J_{\theta}(q_{t+1}^*)$$

This step involves maximizing the expected log-likelihood of the complete data (observed and latent), which is typically solved by differentiating the objective with respect to  $\theta$  and setting the result to zero.

The following diagram illustrates the iterative nature of the EM algorithm.



*This diagram shows the alternating E-step and M-step of the EM algorithm, which continues until the parameters converge.*

## Convergence Guarantee

(18:16) A key property of the EM algorithm is that it guarantees that the log-likelihood will **never decrease** at each iteration:

$$L(\theta_{t+1}) \geq L(\theta_t)$$

While this does not guarantee convergence to the global maximum, it ensures monotonic improvement, typically leading to a good local maximum.

---

## Application: Gaussian Mixture Models (GMM)

(06:55) The instructor uses the Gaussian Mixture Model (GMM) as a concrete example from classical machine learning to illustrate the EM algorithm.

### GMM as a Latent Variable Model

A GMM is a probabilistic model that assumes the observed data is generated from a mixture of several Gaussian distributions.

- **Latent Variable  $z$ :** (07:31) For a GMM with  $M$  components, the latent variable  $z$  is discrete and indicates which of the  $M$  Gaussians generated a data point  $x$ .

$$z \in \{1, 2, \dots, M\}$$

- **Model Parameters  $\theta$ :** (10:25) The parameters of a GMM are:
  - **Mixing Coefficients  $\alpha_j$ :** The prior probability of selecting component  $j$ , where  $p_\theta(z = j) = \alpha_j$ . These must satisfy  $\alpha_j \geq 0$  and  $\sum_{j=1}^M \alpha_j = 1$ .
  - **Component Means  $\mu_j$ :** The mean of each Gaussian component,  $\mu_j \in \mathbb{R}^d$ .
  - **Component Covariances  $\Sigma_j$ :** The covariance matrix of each Gaussian component,  $\Sigma_j \in \mathbb{R}^{d \times d}$ . The full set of parameters is  $\theta = \{\alpha_j, \mu_j, \Sigma_j\}_{j=1}^M$ .
- **Likelihood Function:** The likelihood of a data point  $x$  is a weighted sum of the Gaussian densities:

$$p_\theta(x) = \sum_{j=1}^M \alpha_j \cdot \mathcal{N}(x; \mu_j, \Sigma_j)$$

## EM Algorithm for GMM

(18:53) For a GMM, the steps of the EM algorithm have a clear, analytical form.

- **E-Step:** (19:06) This step corresponds to computing the posterior probability of the latent variable  $z$  given the data  $x$  and the current parameters  $\theta_t$ . This is often called the “responsibility” that component  $j$  takes for data point  $x$ .

$$q_{t+1}^*(z = j|x) = p_{\theta_t}(z = j|x) = \frac{p_{\theta_t}(x|z = j)p_{\theta_t}(z = j)}{\sum_{k=1}^M p_{\theta_t}(x|z = k)p_{\theta_t}(z = k)} = \frac{\alpha_j^t \mathcal{N}(x; \mu_j^t, \Sigma_j^t)}{\sum_{k=1}^M \alpha_k^t \mathcal{N}(x; \mu_k^t, \Sigma_k^t)}$$

- **M-Step:** (22:53) This step involves updating the parameters  $\theta$  to maximize the ELBO, using the responsibilities computed in the E-step. This leads to closed-form update rules for  $\alpha_j$ ,  $\mu_j$ , and  $\Sigma_j$ .

## Limitations and The Path to Deep Generative Models

(26:00) The instructor concludes by highlighting a critical limitation of the standard EM algorithm, which motivates the development of more powerful generative models.

- **The Intractability Problem Revisited:** The EM algorithm is only feasible if the posterior distribution  $p_\theta(z|x)$  can be computed. This is true for GMMs but is **not true** for more complex models where  $p_\theta(x)$  is defined by an intricate function, such as a deep neural network. In such cases, the denominator  $p_\theta(x) = \int p_\theta(x, z) dz$  is intractable, making the posterior  $p_\theta(z|x) = \frac{p_\theta(x, z)}{p_\theta(x)}$  also intractable.
- **The Core Question for Modern Generative Models:** (28:05) > How do we learn a latent variable model for cases where the posterior  $p_\theta(z|x)$  is unknown or intractable?

This question is the central challenge that models like **Variational Autoencoders (VAEs)** and **Diffusion Models** are designed to solve. They use neural networks to approximate these intractable distributions, extending the fundamental principles of ELBO and EM into the deep learning era.

## Self-Assessment for This Video

1. **Explain** why the log-likelihood function  $L(\theta) = \log \int p_\theta(x, z) dz$  is generally intractable for latent variable models.

2. **Derive** the Evidence Lower Bound (ELBO) from the log-likelihood, clearly stating where and why Jensen's Inequality is applied.
3. The ELBO is a function of which two quantities? Write down the final joint optimization problem that we aim to solve.
4. Describe the two alternating steps of the Expectation-Maximization (EM) algorithm. What is being optimized in each step, and what is held constant?
5. In the context of a GMM, what does the latent variable  $z$  represent? What are the three types of parameters that constitute  $\theta$ ?
6. What is the key condition that must be met for the standard EM algorithm to be applicable? Why does this condition fail for complex models like VAEs?

## Key Takeaways from This Video

- **ELBO is a Tractable Proxy:** Direct optimization of log-likelihood in latent variable models is often impossible. The ELBO provides a computable lower bound that we can maximize instead.
- **Learning is Joint Optimization:** Training a latent variable model requires finding the best model parameters ( $\theta$ ) and the best approximation for the posterior ( $q$ ) simultaneously.
- **EM is an Iterative Solution:** The EM algorithm provides an elegant, iterative way to perform this joint optimization by alternating between updating the posterior approximation (E-step) and updating the model parameters (M-step).
- **GMM is a Foundational Example:** Gaussian Mixture Models are a classic, analytically tractable case where the EM algorithm can be applied directly.
- **Intractable Posteriors Motivate Deep Generative Models:** The inability to compute the true posterior  $p_\theta(z|x)$  in complex models is the primary reason for the development of advanced methods like VAEs, which use neural networks to learn an approximation.