

# Study Material - Youtube

## Document Information

- **Generated:** 2025-08-26 06:58:42
- **Source:** <https://www.youtube.com/watch?v=0p7T-3WiPnQ>
- **Platform:** Youtube
- **Word Count:** 2,062 words
- **Estimated Reading Time:** ~10 minutes
- **Number of Chapters:** 1
- **Transcript Available:** Yes (analyzed from video content)

## Table of Contents

1. Inference and Sampling in DDPMs
- 

## Video Overview

This lecture provides a comprehensive explanation of the **inference or sampling process** in Denoising Diffusion Probabilistic Models (DDPMs). The instructor details how to generate new data samples, such as images, by reversing the diffusion (noising) process. Unlike models like GANs or VAEs that generate samples in a single forward pass, DDPMs employ an iterative, multi-step procedure. The lecture breaks down the theoretical foundation of this reverse process, presents the key mathematical formulas, and outlines the complete sampling algorithm.

## Learning Objectives

Upon completing this study material, students will be able to: - **Understand the conceptual difference** between sampling in DDPMs and other generative models like GANs and VAEs. - **Describe the iterative nature** of the DDPM inference process as a reversal of the forward (noising) process. - **Explain the mathematical formulation** of the reverse (decoding) process, including the role of Gaussian transitions. - **Derive and interpret the formulas** for the mean and variance of the reverse conditional distributions. - **Outline the complete step-by-step algorithm** for generating a sample from a trained DDPM, starting from pure noise. - **Recognize the role of the trained neural network** (e.g., U-Net) in predicting the necessary parameters for each denoising step. - **Appreciate the trade-offs** of DDPMs, particularly the high-quality sample generation versus the computationally intensive, slow inference speed.

## Prerequisites

To fully grasp the concepts in this lecture, students should have a foundational understanding of: - The **forward (diffusion) process** in DDPMs, where noise is incrementally added to data. - The **training objective of DDPMs**, including the concept of the Evidence Lower Bound (ELBO) and the simplified loss function. - **Probability theory**, including Gaussian distributions, conditional probability, and Markov chains. - The **reparameterization trick** for sampling from a distribution. - Basic knowledge of neural networks and their role in function approximation.

## Key Concepts Covered in This Video

- **Inference vs. Sampling:** These terms are used interchangeably to describe the process of generating new data.
- **Backward/Decoding Process:** The iterative procedure of removing noise to transform a random noise vector into a coherent data sample.
- **Iterative Sampling:** The core mechanism of DDPM inference, involving multiple steps to generate a single sample.

- **Gaussian Transitions:** The assumption that each step in the reverse process is a Gaussian distribution.
- **Reparameterization Trick:** A technique used to sample from the learned Gaussian distributions in a differentiable way.
- **Trained Neural Network (U-Net):** The model that predicts the parameters (mean or noise) for each step of the reverse process.

## Inference and Sampling in DDPMs

### 1. Intuitive Foundation: From Noise to Data

The primary goal of a generative model is to create new data samples that resemble a training dataset. In the context of DDPMs, this generation process is called **inference** or **sampling**.

At a high level, the process is the exact reverse of the diffusion (or forward) process used during training.

- **Forward Process (Noising):** Starts with a clean image ( $x_0$ ) and iteratively adds a small amount of Gaussian noise at each step to produce a sequence of increasingly noisy images:  $x_1, x_2, \dots, x_T$ . After a large number of steps ( $T$ ), the image  $x_T$  is indistinguishable from pure Gaussian noise. - **Reverse Process (Denoising/Inference):** Starts with a sample of pure Gaussian noise ( $x_T$ ) and iteratively *removes* a small amount of noise at each step, guided by a trained neural network. This generates a sequence that moves from noise to a structured image:  $x_T \rightarrow x_{T-1} \rightarrow \dots \rightarrow x_0$ . The final output,  $x_0$ , is the newly generated sample.

**Key Distinction from GANs/VAEs (00:34):** Unlike Generative Adversarial Networks (GANs) or Variational Autoencoders (VAEs), where a single random vector is passed through a generator network to produce a sample in one shot, DDPMs require an iterative process. This multi-step “denoising” is fundamental to how DDPMs achieve high-quality samples but also explains their slower inference speed.

The entire inference process can be visualized as traversing a Markov chain backward, from the final noisy state to the initial clean state.

flowchart LR

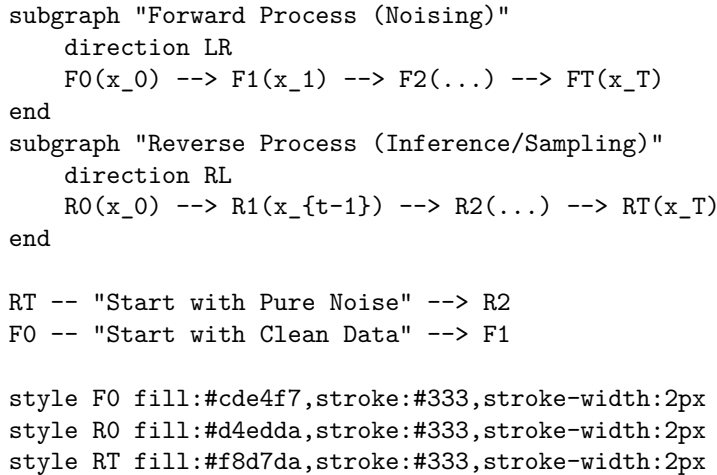


Figure 1: A conceptual diagram illustrating the forward (noising) and reverse (inference) processes in a DDPM. Inference starts with pure noise ( $x_T$ ) and iteratively denoises it to produce a clean sample ( $x_0$ ).

### 2. Mathematical Analysis of the Reverse Process

To generate a sample  $x_0$  from the true data distribution  $p(x_0)$ , we must traverse the reverse Markov chain (01:00). This means we need to iteratively sample from the conditional distributions that define this chain.

The process is as follows: 1. **Start at step T:** Sample  $x_T$  from a standard isotropic Gaussian distribution. This is our initial canvas of pure noise.

$$x_T \sim \mathcal{N}(0, I)$$

2. **Iterate backwards:** For each timestep  $t$  from  $T$  down to 1, we sample  $x_{t-1}$  from the conditional distribution  $p_\theta(x_{t-1}|x_t)$ .

$$x_{t-1} \sim p_\theta(x_{t-1}|x_t) \quad \text{for } t = T, T-1, \dots, 1$$

The final result of this iterative process is  $x_0$ , our generated sample.

## 2.1. The Reverse Transition Distribution

As established during the derivation of the training objective, the reverse transition distribution  $p_\theta(x_{t-1}|x_t)$  is parameterized as a Gaussian (03:53):

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \sigma_q^2(t)I)$$

- **Intuitive Meaning:** This formula states that to get the slightly less noisy image  $x_{t-1}$  from the current noisy image  $x_t$ , we sample from a Gaussian distribution.
- **Parameters:**
  - **Mean  $\mu_\theta(x_t, t)$ :** The center of the distribution for  $x_{t-1}$ . This is the most likely “denoised” version of  $x_t$ . This mean is predicted by our trained neural network, which is why it is parameterized by  $\theta$ .
  - **Variance  $\sigma_q^2(t)I$ :** The variance of the denoising step. In the original DDPM paper, this is a fixed, non-learned hyperparameter, often set to  $\beta_t$  or a related value  $\tilde{\beta}_t$ . It controls how much randomness is injected at each denoising step.

## 2.2. Computing the Mean $\mu_\theta(x_t, t)$

The neural network (typically a U-Net) is trained to predict the noise  $\epsilon_t$  that was added to get from  $x_0$  to  $x_t$ . However, we can use this predicted noise to derive an estimate for the original clean image,  $\hat{x}_0$ , and from there, the mean of the reverse process,  $\mu_\theta(x_t, t)$ .

The formula for the mean of the posterior  $q(x_{t-1}|x_t, x_0)$  was previously derived as:

$$\tilde{\mu}_t(x_t, x_0) = \frac{\sqrt{\bar{\alpha}_{t-1}}\beta_t}{1 - \bar{\alpha}_t}x_0 + \frac{\sqrt{\bar{\alpha}_t}(1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}x_t$$

During training, we train a neural network  $\epsilon_\theta(x_t, t)$  to approximate the true noise  $\epsilon_t$ . We can then get an estimate of  $x_0$  from  $x_t$  using the relation  $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon_t$ :

$$\hat{x}_0(x_t, t) = \frac{1}{\sqrt{\bar{\alpha}_t}}(x_t - \sqrt{1 - \bar{\alpha}_t}\epsilon_\theta(x_t, t))$$

By substituting this  $\hat{x}_0$  into the equation for  $\tilde{\mu}_t$ , we get the formula for the mean of our learned reverse process,  $\mu_\theta(x_t, t)$ :

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\bar{\alpha}_t}} \left( x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right)$$

This is the mean we use for sampling at each step.

### 2.3. Sampling with the Reparameterization Trick

To sample  $x_{t-1}$  from the distribution  $\mathcal{N}(\mu_\theta(x_t, t), \sigma_t^2 I)$ , we use the reparameterization trick (04:22):

1. Calculate the mean  $\mu_\theta(x_t, t)$  using the formula above.
2. Sample a random noise vector  $\epsilon$  from a standard normal distribution:  $\epsilon \sim \mathcal{N}(0, I)$ .
3. Compute the new sample  $x_{t-1}$ :

$$x_{t-1} = \mu_\theta(x_t, t) + \sigma_t \epsilon$$

where  $\sigma_t$  is the standard deviation of the reverse process. For  $t = 1$ , we set  $\epsilon = 0$  to make the final step deterministic.

### 3. The Complete DDPM Inference Algorithm

The entire sampling process can be summarized into a concrete algorithm.

flowchart TD

```

A["Start: Sample noise  
$x_T \sim \mathcal{N}(0, I)$"] --> B["Loop for $t = T$ down to $1$"]
B --> C["Sample noise if $t > 1$  
$\epsilon \sim \mathcal{N}(0, I)$"]
C --> D["Predict noise with NN  
$\epsilon_\theta = \text{NN}(x_t, t)$"]
D --> E["Calculate mean  
$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}}(x_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}}\hat{\epsilon}_t)$"]
E --> F["Sample previous step  
$x_{t-1} = \mu_\theta(x_t, t) + \sigma_t \epsilon$"]
F --> B
B -->|Loop Finished| G["End: Return $x_0$"]

subgraph "For each timestep t"
    C
    D
    E
    F
end
end

```

Figure 2: Flowchart of the DDPM inference algorithm. The process iteratively denoises a random vector over  $T$  steps to generate a final sample.

#### Algorithm: DDPM Sampling

##### 1. Initialization:

- Start with a random sample from a standard Gaussian distribution:

$$x_T \sim \mathcal{N}(0, I)$$

##### 2. Iterative Denoising Loop (06:56):

- For  $t = T, T-1, \dots, 1$ :
  - a. Sample a noise vector  $\epsilon \sim \mathcal{N}(0, I)$ . If  $t = 1$ , set  $\epsilon = 0$ .
  - b. Use the trained neural network  $\epsilon_\theta$  to predict the noise component from the current noisy image  $x_t$ :

$$\hat{\epsilon}_t = \epsilon_\theta(x_t, t)$$

- c. Calculate the mean of the reverse distribution:

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}} \hat{\epsilon}_t \right)$$

- d. Calculate the denoised sample for the previous step:

$$x_{t-1} = \mu_\theta(x_t, t) + \sigma_t \epsilon$$

(where  $\sigma_t^2$  is the fixed variance, e.g.,  $\beta_t$ ).

### 3. Output:

- Return  $x_0$ . This is the final generated sample.

**Important Note on Computational Cost (10:01):** To generate a single sample, this algorithm requires **T forward passes** through the neural network. Since  $T$  is typically large (e.g., 1000), this makes DDPM inference significantly slower than single-pass models like GANs. This is a primary trade-off for the high sample quality and stable training offered by DDPMs.

## Visual Elements from the Video

The lecture heavily relies on handwritten equations and diagrams to explain the concepts.

- **Reverse Process Chain (01:25):** The instructor writes out the sequence  $X_T \rightarrow X_{T-1} \rightarrow X_{T-2} \rightarrow \dots \rightarrow X_0$  to emphasize the step-by-step nature of the decoding process.
- **Iterative Sampling Loop (02:28):** The process of obtaining each state from the previous one is shown as a chain of sampling operations:
  - $X_T \sim \mathcal{N}(0, I)$
  - $X_{T-1} \sim p_\theta(X_{T-1}|X_T)$
  - $X_{T-2} \sim p_\theta(X_{T-2}|X_{T-1})$
  - ...
  - $X_0 \sim p_\theta(X_0|X_1)$
- **U-Net Forward Pass for Inference (07:52):** A diagram shows the trained U-Net taking  $x_t$  and  $t$  as input and producing an estimate of the original clean image,  $\hat{x}_0(x_t)$ . This is described as a “forward pass to get  $\hat{x}_0(x_t)$ ”. This highlights that inference uses the same trained network but in a different, iterative manner.

## Self-Assessment for This Video

1. **Question:** How does the inference process in a DDPM differ from that of a GAN? Why is DDPM inference generally slower?
2. **Question:** What is the starting point of the DDPM sampling process, and what distribution is it drawn from?
3. **Question:** Write the general form of the reverse transition probability  $p_\theta(x_{t-1}|x_t)$ . What are the key parameters of this distribution, and how are they determined during inference?
4. **Problem:** Given a noisy image  $x_t$ , the predicted noise  $\epsilon_\theta(x_t, t)$ , and the noise schedule parameters  $\alpha_t$  and  $\bar{\alpha}_t$ , write down the full expression for sampling  $x_{t-1}$ .
5. **Conceptual Question:** The instructor mentions that the reverse process can be seen as “denoising.” Explain this analogy. What is being “denoised” at each step?
6. **Application Question:** If you have a trained DDPM with  $T = 1000$ , how many times must you run the neural network to generate one new image?

## Key Takeaways from This Video

- **Inference is Iterative Denoising:** Generating a sample in a DDPM is a multi-step process that starts with pure noise and gradually refines it into a structured data point by reversing the diffusion process.
- **The Reverse Process is a Learned Markov Chain:** Each step in the reverse process,  $p_\theta(x_{t-1}|x_t)$ , is a Gaussian distribution whose parameters are determined by the trained neural network.
- **The Neural Network is a Noise Predictor:** During inference, the trained U-Net takes a noisy image  $x_t$  and timestep  $t$  to predict the noise that was added, which in turn allows us to calculate the mean of the next, less-noisy state  $x_{t-1}$ .
- **Inference is Computationally Expensive:** Generating a single sample requires  $T$  sequential forward passes through the neural network, making it much slower than single-shot generative models.
- **Training is Stable, Inference is Slow:** DDPMs are known for their stable training dynamics (avoiding issues like mode collapse in GANs), but this comes at the cost of slow sample generation.